

# Using Python with R

Daniel Chen @chendaniely


DCR Conference 2019

hi!

# I'm Daniel



- PhD Student: Virginia Tech
  - Data Science education
  - Medical practitioners
- Intern at RStudio
  - `gradethis`
  - Code grader for `learnr` documents
- Author:



# R and Python

## The Tiobe Index Top 10

Following are the top 10 languages in the June 2019 Tiobe index:

1. Java
2. C
3. **Python**
4. C
5. Visual
6. C
7. JavaScript
8. PHP
9. SQL
10. Assembly

## The Pypl Index Top 10

Following are the top 10 languages in the June 2018 Pypl index:

- **Python**
- Java
- JavaScript
- C
- PHP
- C
- **R**
- Objective
- Swift
- Matlab

# Python...

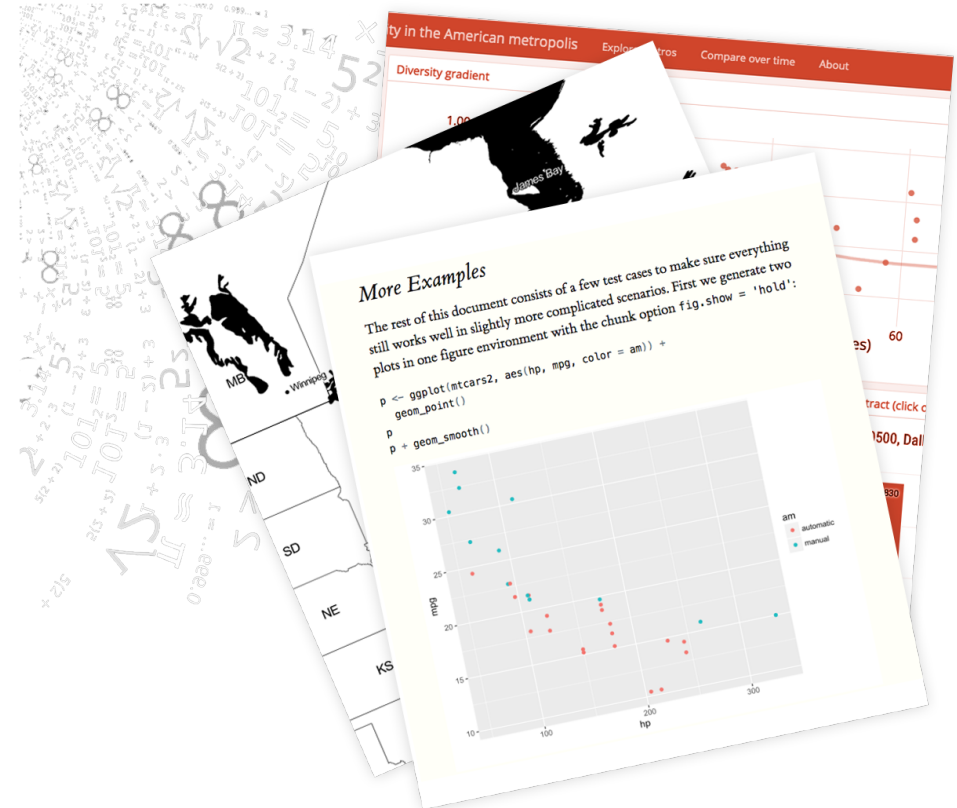
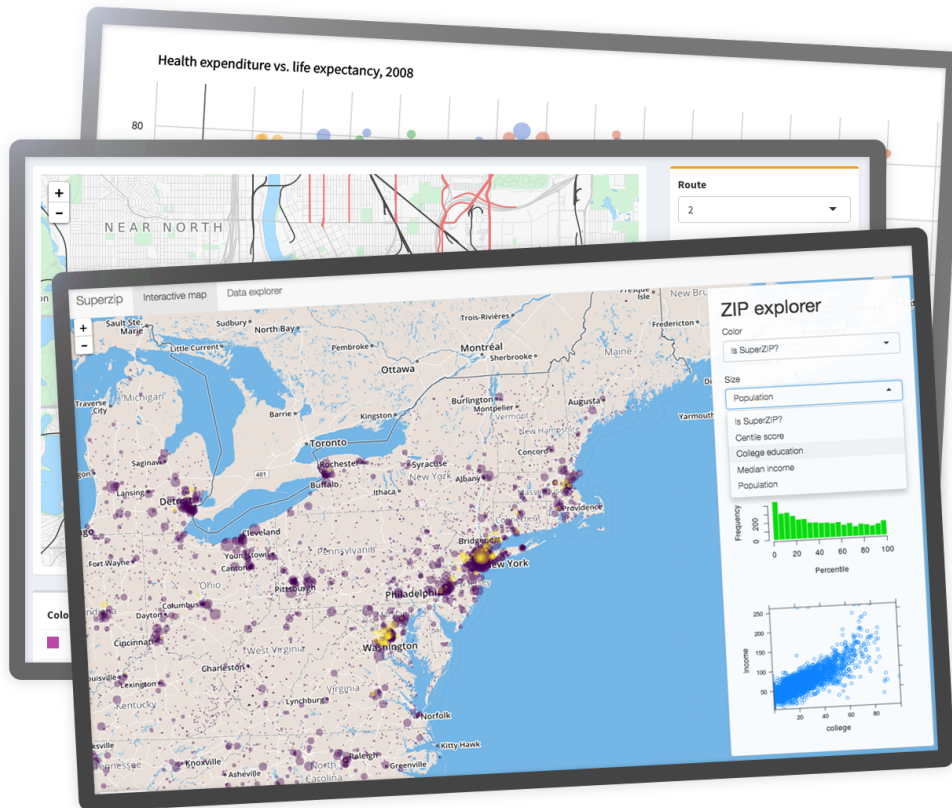
... a general-purpose programming language.

- May not be the best at everything (anything?)
  - but second best at everything is pretty good.
- Python does environments better than R (waiting to test out `renv`)
- One thing that Python is objectively better at than R is *Web Development* and *Hardware*



# What I like about R

## Communication



# Inspiration for talk

2019 Nonclinical Biostatistics Conference

- <https://github.com/chendaniely/ncb-2019-python>
- Jupyter notebook
  - RISE plugin (reveal.js)
  - Slow and clunky
  - Unable to see source (nicely) without Jupyter loaded
- RMarkdown + Reticulate = Slides! (hint hint: this talk ;D)

# R and Python



# R analysis - Load data

```
library(here)
library(readr)

raw = readr::read_csv(here::here("./data/billboard.csv"))
head(raw)
```

```
## # A tibble: 6 x 83
##   year artist.inverted track time genre date.entered date.peaked
##   <dbl> <chr>          <chr> <tim> <chr> <date>      <date>
## 1  2000 Destiny's Child Inde... 03:38 Rock  2000-09-23  2000-11-18
## 2  2000 Santana          Mari... 04:18 Rock  2000-02-12  2000-04-08
## 3  2000 Savage Garden    I Kn... 04:07 Rock  1999-10-23  2000-01-29
## 4  2000 Madonna          Music 03:45 Rock  2000-08-12  2000-09-16
## 5  2000 Aguilera, Chri... Come... 03:38 Rock  2000-08-05  2000-10-14
## 6  2000 Janet            Does... 04:17 Rock  2000-06-17  2000-08-26
## # ... with 76 more variables: x1st.week <dbl>, x2nd.week <dbl>,
## #   x3rd.week <dbl>, x4th.week <dbl>, x5th.week <dbl>, x6th.week <dbl>,
## #   x7th.week <dbl>, x8th.week <dbl>, x9th.week <dbl>, x10th.week <dbl>,
## #   x11th.week <dbl>, x12th.week <dbl>, x13th.week <dbl>,
## #   x14th.week <dbl>, x15th.week <dbl>, x16th.week <dbl>,
## #   x17th.week <dbl>, x18th.week <dbl>, x19th.week <dbl>,
## #   x20th.week <dbl>, x21st.week <dbl>, x22nd.week <dbl>,
## #   x23rd.week <dbl>, x24th.week <dbl>, x25th.week <dbl>,
## #   x26th.week <dbl>, x27th.week <dbl>, x28th.week <dbl>,
## #   x29th.week <dbl>, x30th.week <dbl>, x31st.week <dbl>
```

# R analysis - Filter data

```
library(dplyr)

raw_filtered <- raw %>%
  dplyr::select(year, artist.inverted, track, time, date.entered,
                x1st.week:x73rd.week) %>%
  dplyr::rename(artist = artist.inverted)
raw_filtered
```

```
## # A tibble: 317 x 78
##   year artist track time date.entered x1st.week x2nd.week x3rd.week
##   <dbl> <chr>  <chr> <tim> <date>         <dbl>         <dbl>         <dbl>
## 1  2000 Desti... Inde... 03:38 2000-09-23         78          63          49
## 2  2000 Santa... Mari... 04:18 2000-02-12         15           8           6
## 3  2000 Savag... I Kn... 04:07 1999-10-23         71          48          43
## 4  2000 Madon... Music 03:45 2000-08-12         41          23          18
## 5  2000 Aguil... Come... 03:38 2000-08-05         57          47          45
## 6  2000 Janet  Does... 04:17 2000-06-17         59          52          43
## 7  2000 Desti... Say ... 04:31 1999-12-25         83          83          44
## 8  2000 Igles... Be W... 03:36 2000-04-01         63          45          34
## 9  2000 Sisqo  Inco... 03:52 2000-06-24         77          66          61
## 10 2000 Lones... Amaz... 04:25 1999-06-05         81          54          44
## # ... with 307 more rows, and 70 more variables: x4th.week <dbl>,
## #   x5th.week <dbl>, x6th.week <dbl>, x7th.week <dbl>, x8th.week <dbl>,
## #   x9th.week <dbl>, x10th.week <dbl>, x11th.week <dbl>, x12th.week <dbl>,
## #   x13th.week <dbl>, x14th.week <dbl>, x15th.week <dbl>
```

# R analysis - Tidy data

```
library(tidyr)

raw_tidy <- raw_filtered %>%
  tidyr::pivot_longer(cols = tidyselect::starts_with('x'),
                      names_to = "week",
                      values_to = "rank")

raw_tidy
```

```
## # A tibble: 23,141 x 7
##   year artist      track      time date.entered week      rank
##   <dbl> <chr>      <chr>      <tim> <date>      <chr>    <dbl>
## 1  2000 Destiny's Ch... Independent Women... 03:38 2000-09-23 x1st.we... 78
## 2  2000 Destiny's Ch... Independent Women... 03:38 2000-09-23 x2nd.we... 63
## 3  2000 Destiny's Ch... Independent Women... 03:38 2000-09-23 x3rd.we... 49
## 4  2000 Destiny's Ch... Independent Women... 03:38 2000-09-23 x4th.we... 33
## 5  2000 Destiny's Ch... Independent Women... 03:38 2000-09-23 x5th.we... 23
## 6  2000 Destiny's Ch... Independent Women... 03:38 2000-09-23 x6th.we... 15
## 7  2000 Destiny's Ch... Independent Women... 03:38 2000-09-23 x7th.we... 7
## 8  2000 Destiny's Ch... Independent Women... 03:38 2000-09-23 x8th.we... 5
## 9  2000 Destiny's Ch... Independent Women... 03:38 2000-09-23 x9th.we... 1
## 10 2000 Destiny's Ch... Independent Women... 03:38 2000-09-23 x10th.w... 1
## # ... with 23,131 more rows
```

# R analysis - Clean data

```
library(purrr)
library(stringr)

billboard_clean <- raw_tidy %>%
  dplyr::mutate(
    week = purrr::map_int(
      week,
      #function(x){as.integer(stringr::str_extract(x, '\\d+'))}
      ~ as.integer(stringr::str_extract(., "\\d+"))
    )
  )
billboard_clean
```

```
## # A tibble: 23,141 x 7
##   year artist      track      time date.entered week rank
##   <dbl> <chr>      <chr>      <tim> <date>      <int> <dbl>
## 1  2000 Destiny's Chi... Independent Women P... 03:38 2000-09-23     1     78
## 2  2000 Destiny's Chi... Independent Women P... 03:38 2000-09-23     2     63
## 3  2000 Destiny's Chi... Independent Women P... 03:38 2000-09-23     3     49
## 4  2000 Destiny's Chi... Independent Women P... 03:38 2000-09-23     4     33
## 5  2000 Destiny's Chi... Independent Women P... 03:38 2000-09-23     5     23
## 6  2000 Destiny's Chi... Independent Women P... 03:38 2000-09-23     6     15
## 7  2000 Destiny's Chi... Independent Women P... 03:38 2000-09-23     7      7
## 8  2000 Destiny's Chi... Independent Women P... 03:38 2000-09-23     8      5
## 9  2000 Destiny's Chi... Independent Women P... 03:38 2000-09-23     9      1
```

# Python analysis


```
import pandas as pd
import re

import janitor
from pyprojroot import here

raw_py = pd.read_csv(here('./data/billboard.csv'),
                     encoding = "ISO-8859-1")

billboard_clean_py = (
    raw_py
    .select_columns(['year', 'artist.inverted', 'track', 'time',
                     'date.entered', 'x*.week'])
    .rename_columns({"artist.inverted": "artist"})
    .melt(id_vars = ['year', 'artist', 'track', 'time',
                     'date.entered'],
          var_name = "week",
          value_name = "rank")
    .transform_column('week',
                      lambda x: int(re.findall(r'\d+', x)[0]))
)
```

# Reticulate -- Python in R!



- Calling Python from R
- Translation between R and Python objects
- Python environments

# Reticulate

```
library(reticulate)
(conda_envs <- reticulate::conda_list())
```

```
##           name                               python
## 1 miniconda3 /home/dchen/miniconda3/bin/python
```

```
# use my default conda environment
conda_envs$name[[1]]
```

```
## [1] "miniconda3"
```

```
env <- conda_envs$name[[1]]
reticulate::use_condaenv(env)
```

```
reticulate::py_config()
```

```
## python:           /home/dchen/miniconda3/bin/python
## libpython:        /home/dchen/miniconda3/lib/libpython3.7m.so
## pythonhome:       /home/dchen/miniconda3:/home/dchen/miniconda3
## version:          3.7.3 | packaged by conda-forge | (default, Jul  1 2019, 21:52:21) [GCC 7.3.0]
## numpy:            /home/dchen/miniconda3/lib/python3.7/site-packages/numpy
## numpy version:    1.17.3
```

# A Python script

```
import pandas as pd
import re

import janitor
from pyprojroot import here

raw_py = pd.read_csv(here('./data/billboard.csv'),
                     encoding = "ISO-8859-1")

billboard_clean_py = (
    raw_py
    .select_columns(['year', 'artist.inverted', 'track', 'time',
                    'date.entered', 'x*.week'])
    .rename_columns({"artist.inverted": "artist"})
    .melt(id_vars = ['year', 'artist', 'track', 'time',
                    'date.entered'],
          var_name = "week",
          value_name = "rank")
    .transform_column('week',
                      lambda x: int(re.findall(r'\d+', x)[0]))
)

mean_rank_by_week = (billboard_clean_py.groupby("week")["rank"]
                      .mean())
```



# Python objects in R

```
reticulate::source_python(here::here("./scripts/01-02-python.py"))
```

```
head(mean_rank_by_week)
```

```
##           1           2           3           4           5           6
## 79.95899 71.17308 65.04560 59.76333 56.33904 52.36071
```

```
head(billboard_clean_py)
```

```
##   year          artist          track time
## 1 2000   Destiny's Child Independent Women Part I 3:38
## 2 2000          Santana          Maria, Maria 4:18
## 3 2000   Savage Garden          I Knew I Loved You 4:07
## 4 2000          Madonna          Music 3:45
## 5 2000 Aguilera, Christina Come On Over Baby (All I Want Is You) 3:38
## 6 2000          Janet          Doesn't Really Matter 4:17
##   date.entered week rank
## 1 2000-09-23     1    78
## 2 2000-02-12     1    15
## 3 1999-10-23     1    71
## 4 2000-08-12     1    41
## 5 2000-08-05     1    57
## 6 2000-06-17     1    59
```

# Type conversions table

| R                      | Python            | Examples  |
|------------------------|-------------------|---|
| Single-element vector  | Scalar            | <code>1, 1L, TRUE, "foo"</code>                             |
| Multi-element vector   | List              | <code>c(1.0, 2.0, 3.0), c(1L, 2L, 3L)</code>                |
| List of multiple types | Tuple             | <code>list(1L, TRUE, "foo")</code>                          |
| Named list             | Dict              | <code>list(a = 1L, b = 2.0), dict(x = x_data)</code>        |
| Matrix/Array           | NumPy ndarray     | <code>matrix(c(1,2,3,4), nrow = 2, ncol = 2)</code>         |
| Data Frame             | Pandas DataFrame  | <code>data.frame(x = c(1,2,3), y = c("a", "b", "c"))</code> |
| Function               | Python function   | <code>function(x) x + 1</code>                              |
| NULL, TRUE, FALSE      | None, True, False | NULL, TRUE, FALSE   |

<https://rstudio.github.io/reticulate/#type-conversions>

# Machine Learning

# The data

No standard for how to transport data within a package...

```
from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer()

print(type(cancer))
```

```
## <class 'sklearn.utils.Bunch'>
```

**vs**

```
import seaborn as sns
tips = sns.load_dataset("tips")
print(type(tips))
```

```
## <class 'pandas.core.frame.DataFrame'>
```

# The data

```
cancer.target[:10]
```

```
## array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
cancer.data[:10]
```

```
## array([[1.799e+01, 1.038e+01, 1.228e+02, 1.001e+03, 1.184e-01, 2.776e-01,
##         3.001e-01, 1.471e-01, 2.419e-01, 7.871e-02, 1.095e+00, 9.053e-01,
##         8.589e+00, 1.534e+02, 6.399e-03, 4.904e-02, 5.373e-02, 1.587e-02,
##         3.003e-02, 6.193e-03, 2.538e+01, 1.733e+01, 1.846e+02, 2.019e+03,
##         1.622e-01, 6.656e-01, 7.119e-01, 2.654e-01, 4.601e-01, 1.189e-01],
##        [2.057e+01, 1.777e+01, 1.329e+02, 1.326e+03, 8.474e-02, 7.864e-02,
##         8.690e-02, 7.017e-02, 1.812e-01, 5.667e-02, 5.435e-01, 7.339e-01,
##         3.398e+00, 7.408e+01, 5.225e-03, 1.308e-02, 1.860e-02, 1.340e-02,
##         1.389e-02, 3.532e-03, 2.499e+01, 2.341e+01, 1.588e+02, 1.956e+03,
##         1.238e-01, 1.866e-01, 2.416e-01, 1.860e-01, 2.750e-01, 8.902e-02],
##        [1.969e+01, 2.125e+01, 1.300e+02, 1.203e+03, 1.096e-01, 1.599e-01,
##         1.974e-01, 1.279e-01, 2.069e-01, 5.999e-02, 7.456e-01, 7.869e-01,
##         4.585e+00, 9.403e+01, 6.150e-03, 4.006e-02, 3.832e-02, 2.058e-02,
##         2.250e-02, 4.571e-03, 2.357e+01, 2.553e+01, 1.525e+02, 1.709e+03,
##         1.444e-01, 4.245e-01, 4.504e-01, 2.430e-01, 3.613e-01, 8.758e-02],
##        [1.142e+01, 2.038e+01, 7.758e+01, 3.861e+02, 1.425e-01, 2.839e-01,
##         2.414e-01, 1.052e-01, 2.597e-01, 9.744e-02, 4.956e-01, 1.156e+00,
##         3.445e+00, 2.723e+01, 9.110e-03, 7.458e-02, 5.661e-02, 1.867e-02]
```

# Python -- Preprocess

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.svm import SVC

# split the data
X_train, X_test, y_train, y_test = train_test_split(
    cancer.data, cancer.target, random_state=0)

# compute minimum and maximum on the training data
scaler = MinMaxScaler().fit(X_train)

# rescale training data
X_train_scaled = scaler.transform(X_train)
```

# Python -- Fit

```
svm = SVC()  
# learn an SVM on the scaled training data  
svm.fit(X_train_scaled, y_train)
```

```
## SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,  
##      decision_function_shape='ovr', degree=3, gamma='auto_deprecated',  
##      kernel='rbf', max_iter=-1, probability=False, random_state=None,  
##      shrinking=True, tol=0.001, verbose=False)
```

# Python -- Evaluate

```
# scale test data and score the scaled data  
X_test_scaled = scaler.transform(X_test)  
svm.score(X_test_scaled, y_test)
```

```
## 0.951048951048951
```

Default scoring metric is accuracy



# R -- Python setup

```
library(reticulate)
# reticulate::use_condaenv("miniconda3")

(conda_envs <- reticulate::conda_list())
```

```
##           name                               python
## 1 miniconda3 /home/dchen/miniconda3/bin/python
```

```
conda_envs$name[[1]]
```

```
## [1] "miniconda3"
```

```
env <- conda_envs$name[[1]]
reticulate::use_condaenv(env)
```

# R -- Get data

```
sklearn_datasets = reticulate::import_from_path("sklearn.datasets")
cancer = sklearn_datasets$load_breast_cancer()
```

```
library(tibble)
cancer_df <- tibble::as_tibble(cancer$data)
names(cancer_df) <- cancer$feature_names
cancer_df$target <- cancer$target
cancer_df
```

```
## # A tibble: 569 x 31
##   `mean radius` `mean texture` `mean perimeter` `mean area`
##   <dbl>         <dbl>         <dbl>         <dbl>
## 1         18.0         10.4         123.         1001
## 2         20.6         17.8         133.         1326
## 3         19.7         21.2         130          1203
## 4         11.4         20.4          77.6         386.
## 5         20.3         14.3         135.         1297
## 6         12.4         15.7          82.6         477.
## 7         18.2         20.0         120.         1040
## 8         13.7         20.8          90.2         578.
## 9          13         21.8          87.5         520.
## 10        12.5         24.0          84.0         476.
## # ... with 559 more rows, and 27 more variables: `mean smoothness` <dbl>,
## #   `mean compactness` <dbl>, `mean concavity` <dbl>, `mean concave
## #   points` <dbl>, `mean symmetry` <dbl>, `mean fractal dimension` <dbl>
```

# R -- Preprocess

```
library(rsample)
library(recipes)

cancer_split <- rsample::initial_split(cancer_df)
cancer_train <- rsample::training(cancer_split)
cancer_test <- rsample::testing(cancer_split)

res <- recipes::recipe(target ~ ., data = cancer_train) %>%
  recipes::step_scale(recipes::all_predictors()) %>%
  recipes::step_num2factor(recipes::all_outcomes())

res_prepd <- res %>% recipes::prep()
res_baked <- res_prepd %>% bake(new_data = cancer_train,
                              composition = "tibble")
res_test <- res_prepd %>% bake(new_data = cancer_test,
                              composition = "tibble")
```

# R -- Fit

<https://tidymodels.github.io/parsnip/articles/articles/Models.html>

```
library(parsnip)

svm <- parsnip::svm_rbf(mode = "classification", cost = 1) %>%
  parsnip::set_engine("kernlab") %>%
  parsnip::fit(target ~ ., data = res_baked)
```

# R -- Evaluate

```
library(yardstick)
predict(svm, res_test) %>%
  dplyr::bind_cols(res_test %>% dplyr::select(target)) %>%
  yardstick::accuracy(truth = target, estimate = .pred_class)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>    <chr>         <dbl>
## 1 accuracy binary         0.972
```

# Communication

# This presentation

- Written in RMarkdown exported as a **xaringan** slide deck
- All the R **and** Python code are live executed by changing the execution engine

```
```${r}
library(reticulate)
reticulate::use_condaenv("Anaconda3")
```

```${python}
from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer()
```
```

# Sharing objects R < -- > Python

- In R chunks, you can access python objects with: `py$`
- In Python chunks, you can access R objects with: `r`.
  - Note the dot in `r`.

In Python chunk:

```
single_obs_py = X_test_scaled[:1, :] # first row of the test data
single_obs_py
```

```
## array([[0.30380046, 0.44854772, 0.30993021, 0.17527041, 0.62962963,
##          0.43668242, 0.33856607, 0.40616302, 0.53333333, 0.49052233,
##          0.10106826, 0.12555836, 0.11006926, 0.04942689, 0.17120785,
##          0.1958559 , 0.08717172, 0.25269937, 0.17111501, 0.10745132,
##          0.301672  , 0.47014925, 0.31321281, 0.16201337, 0.56943802,
##          0.34763416, 0.40782748, 0.70651051, 0.39818648, 0.36639118]])
```

In R chunk:

```
single_obs_r <-py$single_obs_py # get an R object
```



# Python -- Prediction

```
svm.predict(single_obs_py) # using python variable
```

```
## array([0])
```

```
svm.predict(r.single_obs_r) # using R variable
```

```
## array([0])
```

# R -- Prediction

```
r_dat <- as.data.frame(single_obs_r)
names(r_dat) <- py$cancer$feature_names
predict(svm, r_dat)
```

```
## # A tibble: 1 x 1
##   .pred_class
##   <fct>
## 1 0
```

# Shiny

[https://scikit-learn.org/stable/modules/model\\_persistence.html](https://scikit-learn.org/stable/modules/model_persistence.html)

Save out the model

```
from joblib import dump, load
from pyprojroot import here

dump(svm, here("output/python_model.joblib", warn=False))
```

```
## ['/home/dchen/git/hub/rstatsdc_2019-python-r/output/python_model.joblib']
```

Load the model

```
python_model = load(here("output/python_model.joblib"))
```

# Shiny

[https://github.com/chendaniely/rstatsdc\\_2019-python-r/blob/master/shiny\\_example.Rmd](https://github.com/chendaniely/rstatsdc_2019-python-r/blob/master/shiny_example.Rmd)

```
```{r}
inputPanel(
  sliderInput("bw_adjust", label = "Bandwidth adjustment:",
             min = 1, max = 20, value = 1, step = 1)
)

renderText({
  py$python_model$predict(py$X_test_scaled[1:input$bw_adjust, , drop=FALSE])
})
```
```

# The -down ecosystem

All of this is using the `reticulate` R package

<https://rstudio.github.io/reticulate/>

- Bookdown
- Blogdown
  - Hugo academic already supports Jupyter notebooks
  - <https://sourcethemes.com/academic/docs/jupyter>

## By the way...

- `knitpy`: <https://github.com/jankatins/knitpy>
- `jupyter` books: <https://jupyterbook.org/intro.html>

# Creating a reticulated R package

The R keras package is an R wrapper around keras for Python

- <https://keras.rstudio.com/>

<https://rstudio.github.io/reticulate/articles/package.html>

# Installing Python...

I recommend looking at the Software-Carpentry setup instructions:

<https://swcarpentry.github.io/python-novice-inflammation/setup/index.html>

Most people in data science use Anaconda to install Python

- <https://www.anaconda.com/distribution/>

People who mainly use python for Web development don't use Anaconda

# About conda..

What they forgot to teach you about R: <https://rstats.wtf/>

There's a section about using conda with R: <https://rstats.wtf/set-up-an-r-dev-environment.html#what-about-conda>

tl;dr - don't mix `conda install` with `install.packages()`



# Apache arrow

If you heard me speak before...

- DCR 2018: Structuring Your Data Science Projects
  - <https://youtu.be/UQHz38s3DyA>
- NYR 2019: Building Reproducible and Replicable Projects
  - <https://youtu.be/t-vY9FeIIMk>

Save out data objects to share between Python and R scripts

- Python: <https://arrow.apache.org/docs/python/>
- R: <https://arrow.apache.org/docs/r/>

# Thanks!

@chendaniely

Slides: [https://github.com/chendaniely/rstatsdc\\_2019-python-r](https://github.com/chendaniely/rstatsdc_2019-python-r)